

HYPER PARAMETER TUNING

```
In [20]: DT = DecisionTreeClassifier ( )

Parameters = { 'criterion': [ 'gini' , 'entropy' ] ,

               'splitter' : [ 'best' , 'random' ] ,

               'max_depth' : [ None , 5 , 10 , 15 ] ,

               'max_features' : [ 'None' , "sqrt" , "log2" ] }

GS = GridSearchCV ( estimator = DT , param_grid = Parameters , cv = 5 )

GS.fit ( Train_X , Train_Y )

Parameters = GS.best_params_

print ( "The best parameters of decision tree are : \n\n" , Parameters )

The best parameters of decision tree are :

{'criterion': 'entropy', 'max_depth': None, 'max_features': 'log2', 'splitter': 'best'}
```

TRAINING AND TESTING

```
In [21]: DT = DecisionTreeClassifier ( **Parameters )

DT.fit ( Train_X , Train_Y )

Y_DT = DT.predict ( Test_X )

Accuracy_DT = np.round ( accuracy_score ( Test_Y , Y_DT ) * 100 , 2 )

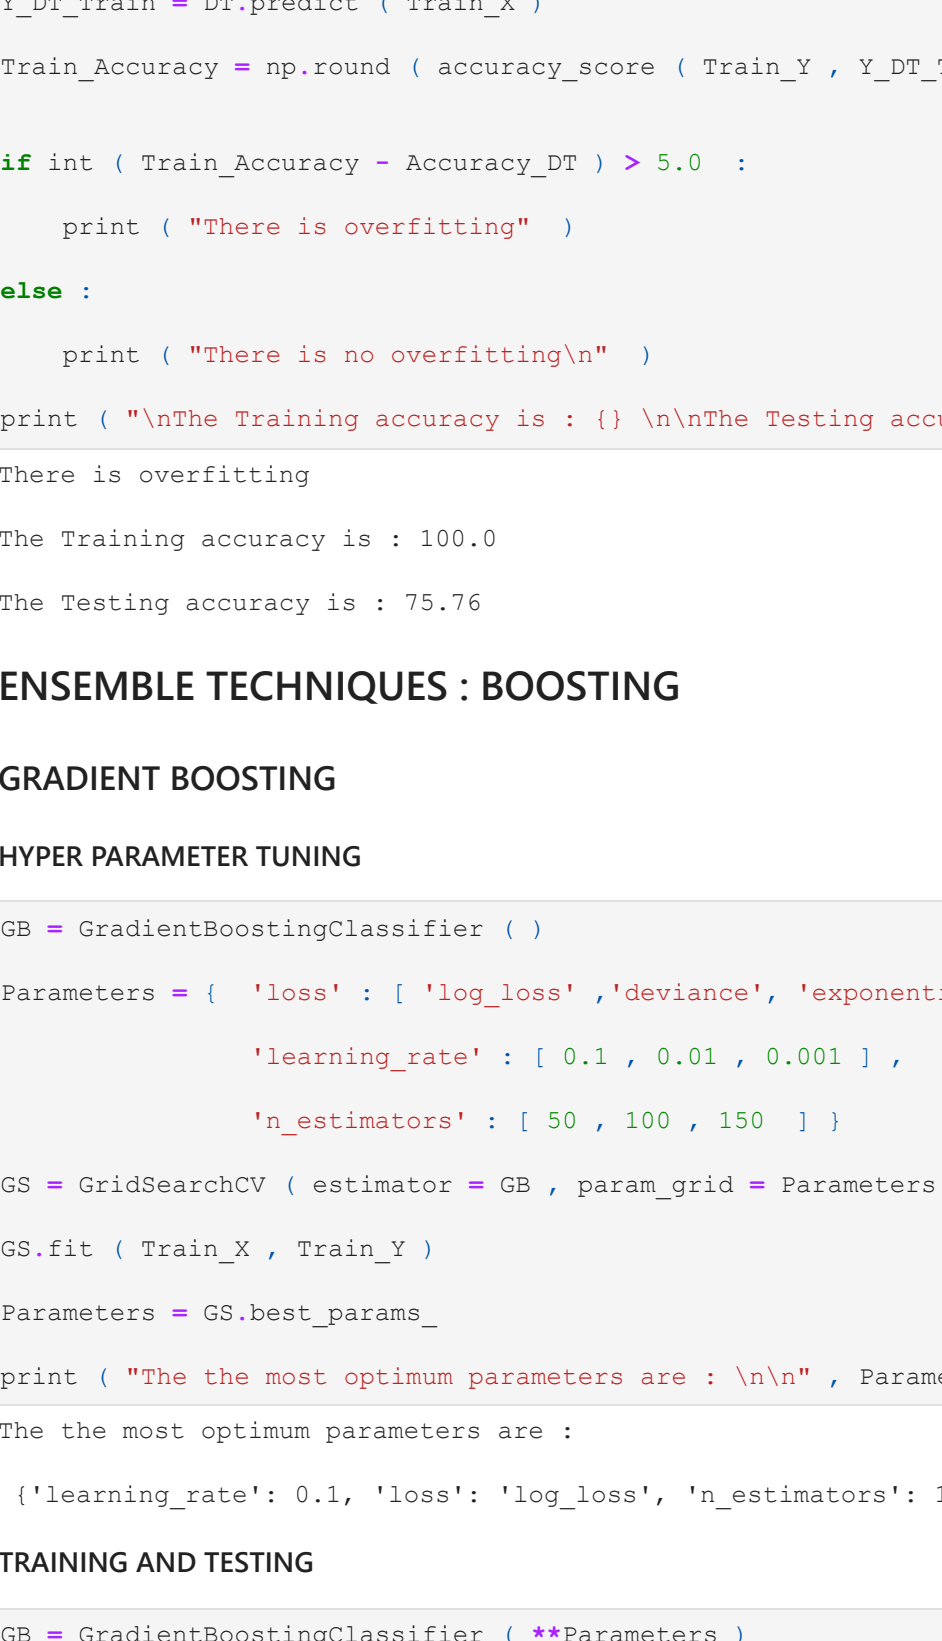
print ( "The accuracy of the test data is : " , Accuracy_DT )

Confusion_Matrix_DT = sns.heatmap ( confusion_matrix ( Test_Y , Y_DT ) , fnt = '0.0f' , annot = True , cmap = '

Confusion_Matrix_DT

The accuracy of the test data is : 75.76

<AxesSubplot>
```



CHECKING FOR OVER FITTING

```
In [22]: Y_DT_Train = DT.predict ( Train_X )

Train_Accuracy = np.round ( accuracy_score ( Train_Y , Y_DT_Train ) * 100 , 2 )

if int ( Train_Accuracy - Accuracy_DT ) > 5.0 :

    print ( "There is overfitting" )

else :

    print ( "There is no overfitting\n" )

print ( "\nThe Training accuracy is : {} \n\nThe Testing accuracy is : {}".format( Train_Accuracy , Accuracy_DT )

There is overfitting

The Training accuracy is : 100.0

The Testing accuracy is : 75.76
```

ENSEMBLE TECHNIQUES : BOOSTING

GRADIENT BOOSTING

HYPER PARAMETER TUNING

```
In [23]: GB = GradientBoostingClassifier ( )

Parameters = ( 'loss': [ 'log_loss' , 'deviance' , 'exponential' ] ,

               'learning_rate' : [ 0.1 , 0.01 , 0.001 ] ,

               'n_estimators' : [ 50 , 100 , 150 ] )

GS = GridSearchCV ( estimator = GB , param_grid = Parameters )

GS.fit ( Train_X , Train_Y )

Parameters = GS.best_params_

print ( "The the most optimum parameters are : \n\n" , Parameters )

The the most optimum parameters are :

{'learning_rate': 0.1, 'loss': 'log_loss', 'n_estimators': 150}
```

TRAINING AND TESTING

```
In [24]: GB = GradientBoostingClassifier ( **Parameters )

GB.fit ( Train_X , Train_Y )

Y_GB = GB.predict ( Test_X )

Accuracy_GB = np.round ( accuracy_score ( Test_Y , Y_GB ) * 100 , 2 )

print ( "Accuracy of gradient boosting is : " , Accuracy_GB )

Confusion_Matrix_GB = sns.heatmap ( confusion_matrix ( Test_Y , Y_GB ) , annot = True , fnt = '0.0f' , cmap = '

Confusion_Matrix_GB

Accuracy of gradient boosting is : 87.88

<AxesSubplot>
```



CHECKING FOR OVER FITTING

```
In [25]: Y_GB_Train = GB.predict ( Train_X )

Train_Accuracy = np.round ( accuracy_score ( Train_Y , Y_GB_Train ) * 100 , 2 )

if int ( Train_Accuracy - Accuracy_GB ) > 5.0 :

    print ( "There is overfitting" )

else :

    print ( "There is no overfitting\n" )

print ( "\nThe Training accuracy is : {} \n\nThe Testing accuracy is : {}".format( Train_Accuracy , Accuracy_GB )

There is overfitting

The Training accuracy is : 100.0

The Testing accuracy is : 87.88
```

EXTREME GRADIENT BOOSTING

HYPER PARAMETER TUNING

```
In [26]: XGB = XGBClassifier ( )

Parameters = {

               'n_estimators': [100, 200, 300],

               'learning_rate': [ 0.5 , 0.1 , 0.01 , 0.001],

               'max_depth': [3, 5, 7]

            }

LE = LabelEncoder ( )

Train_Y = LE.fit_transform ( Train_Y )

Test_Y = LE.fit_transform ( Test_Y )

GS = GridSearchCV ( estimator = XGBClassifier ( ) , param_grid = Parameters )

GS.fit ( Train_X , Train_Y )

Parameters = GS.best_params_

print ( "The the most optimum parameters are : \n\n" , Parameters )

The the most optimum parameters are :

{'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 200}
```

TRAINING AND TESTING

```
In [27]: XGB = XGBClassifier ( **Parameters )

XGB.fit ( Train_X , Train_Y )

Y_XGB = XGB.predict ( Test_X )

Accuracy_XGB = np.round ( accuracy_score ( Test_Y , Y_XGB ) * 100 , 2 )

print ( "Accuracy of extreme gradient boosting is : " , Accuracy_XGB )

Confusion_Matrix_XGB = sns.heatmap ( confusion_matrix ( Test_Y , Y_XGB ) , annot = True , fnt = '0.0f' , cmap = '

Confusion_Matrix_XGB

Accuracy of extreme gradient boosting is : 93.94

<AxesSubplot>
```



CHECKING FOR OVER FITTING

```
In [28]: Y_XGB_Train = XGB.predict ( Train_X )

Train_Accuracy = np.round ( accuracy_score ( Train_Y , Y_XGB_Train ) * 100 , 2 )

if int ( Train_Accuracy - Accuracy_XGB ) > 5.0 :

    print ( "There is overfitting" )

else :

    print ( "There is no overfitting\n" )

print ( "\nThe Training accuracy is : {} \n\nThe Testing accuracy is : {}".format( Train_Accuracy , Accuracy_XGB )

There is overfitting

The Training accuracy is : 100.0

The Testing accuracy is : 93.94
```

COMPARISON

```
In [29]: Accuracies = [ Accuracy_DT , Accuracy_LR , Accuracy_GB , Accuracy_RF , Accuracy_XGB ]

Accuracies_Names = [ 'Accuracy_DT' , 'Accuracy_LR' , 'Accuracy_GB' , 'Accuracy_RF' , 'Accuracy_XGB' ]

Performance_Metrics = { 'Accuracies' : Accuracies , 'Values' : Accuracies_Names }

Performance_Metrics = pd.DataFrame(Performance_Metrics)

Out[29]:
```

| | Accuracies | Values |
|---|------------|--------------|
| 0 | 75.76 | Accuracy_DT |
| 1 | 63.64 | Accuracy_LR |
| 2 | 87.88 | Accuracy_GB |
| 3 | 87.88 | Accuracy_RF |
| 4 | 93.94 | Accuracy_XGB |

CONCLUSION

The results clearly demonstrate that the XGBoost algorithm outperforms the other methods, achieving the highest accuracy among all evaluated classifiers.

Random Forest also performed admirably, showing competitive performance with an accuracy close to XGBoost.

However, Decision Tree and Logistic Regression exhibited relatively lower accuracies, indicating that they might not be the best choices for this specific dataset.